

COMP1204: Modelling and SQL

Oli Bills

ofb@ecs.soton.ac.uk

What we have covered so far

- Database Design Theory
- Functional dependencies
- Keys/Superkeys
- Normalisation
- Next... data modelling

Terminology

- Student as an **entity** (an object, a concept) has a number of **attributes** (ID,Name)
- Student as a **Relation**: Student(ID,Name) has the attributes/columns (ID,Name)
- Student **HAS** a Tutor (that's a **relationship**)

Data Model

- “Things” of importance to the system
- How they relate to each other
- Built / Modified - iterated until suited to system model
- Can be represented as UML class diagrams
 - add behaviour of each entity to the model

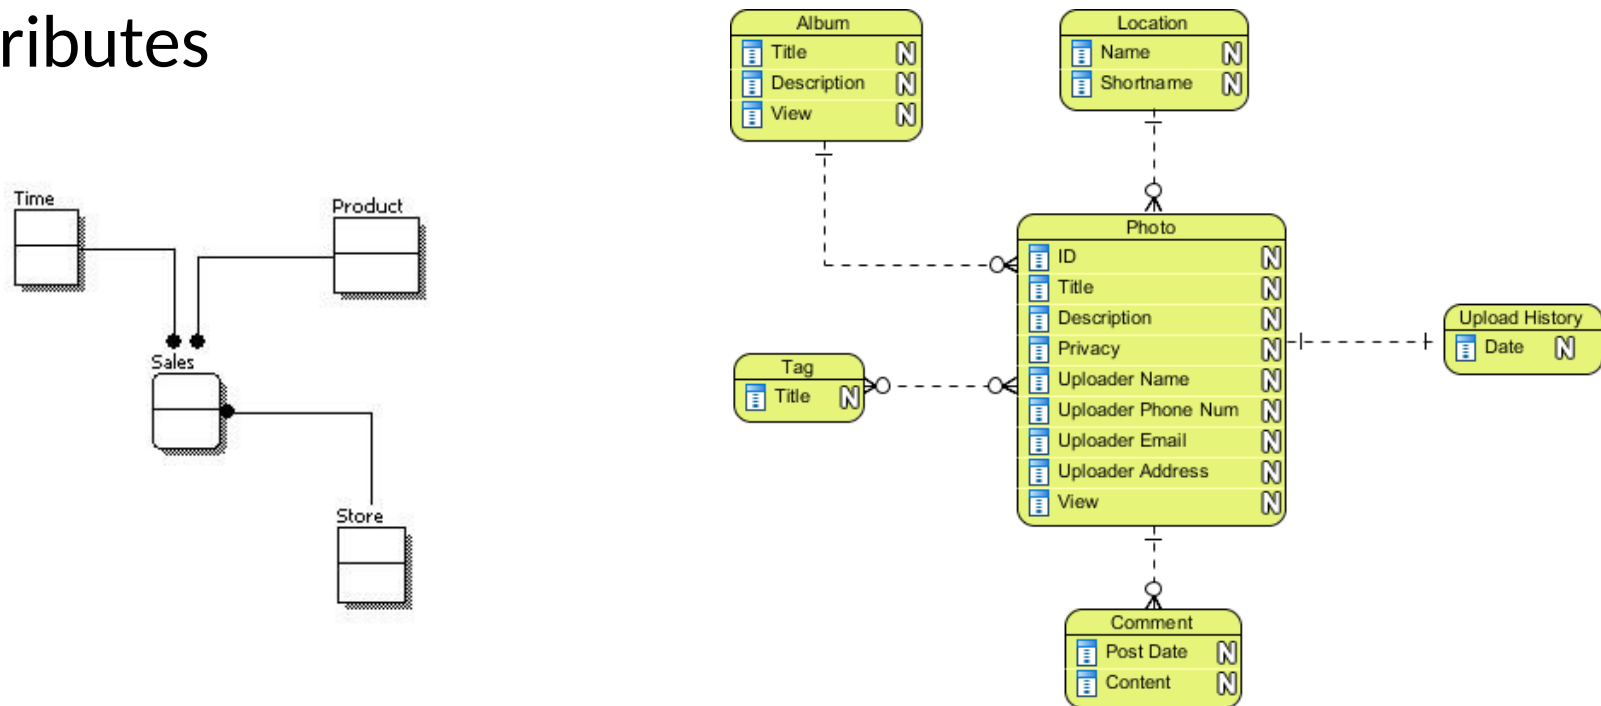
Three types of Data Modelling

- Conceptual -> Ideas
- Logical -> High Level Design
- Physical -> Low Level

Feature	Conceptual	Logical	Physical
Entity names	X	X	
Entity relationships	X	X	
Attributes		X	
Primary keys		X	X
Foreign keys		X	X
Table names			X
Column names			X
Column data types			X

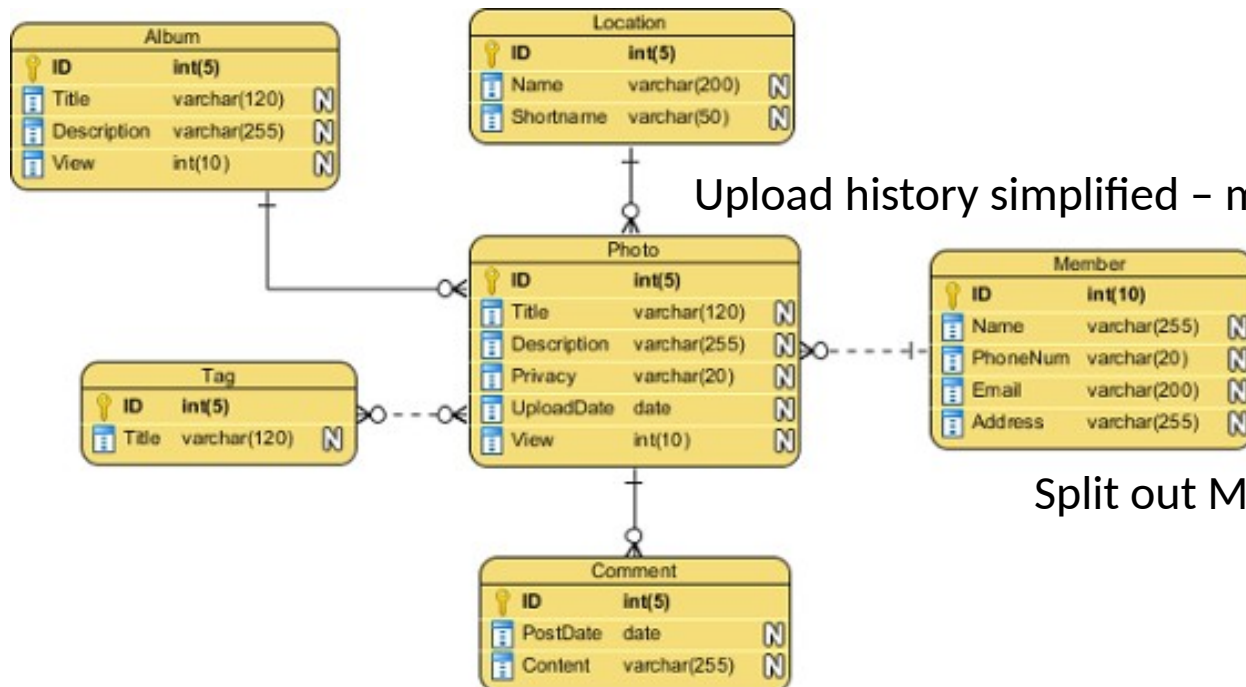
Conceptual Model

- Directly from the requirements and domain
- No thought of database design
- Entity names and relationships and sometimes attributes



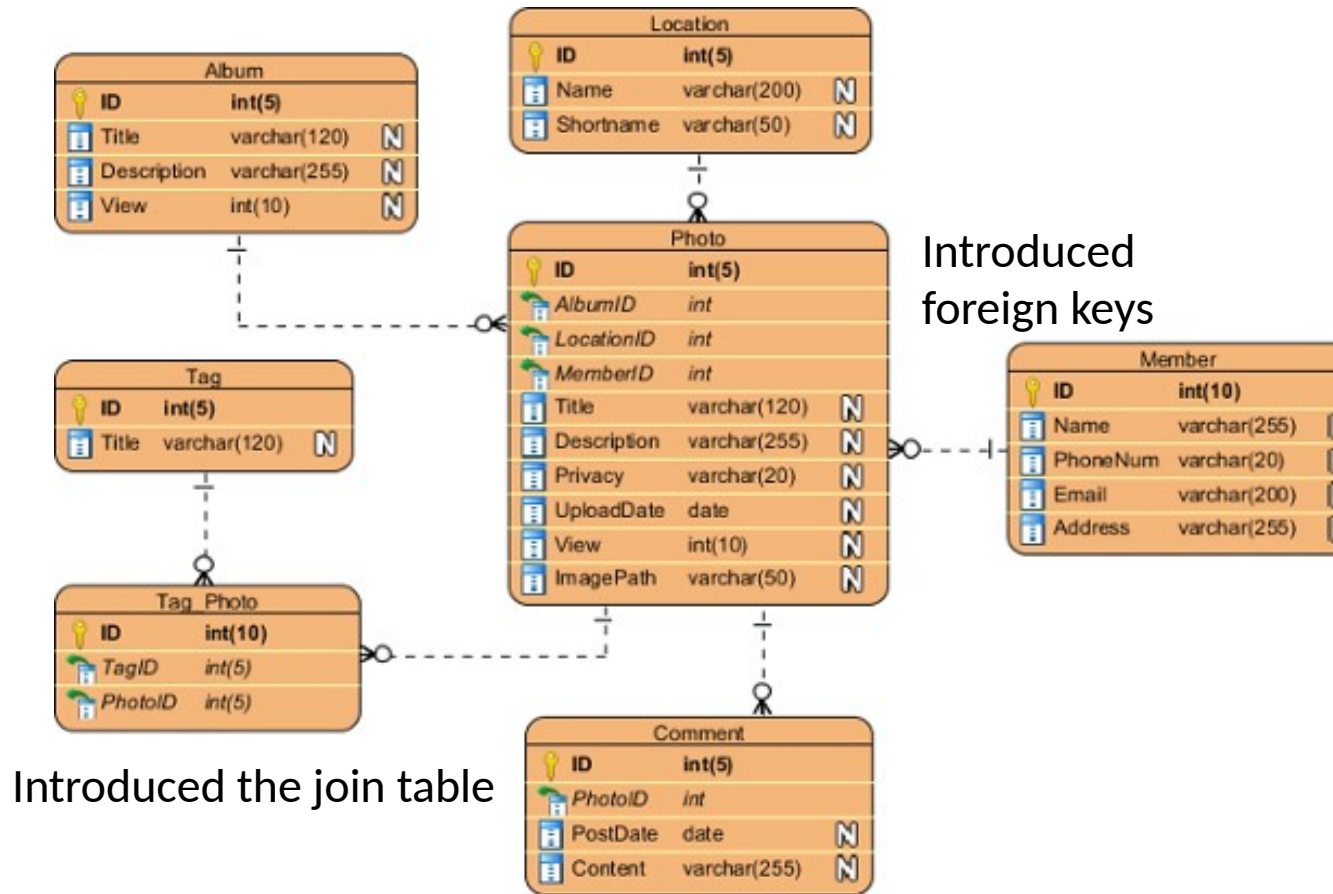
Logical Model

- Identification of attributes
- Identification of attribute types



Physical Model

- Consideration of database structure
- Actual tables and fields
- Implementation of relationships (e.g. keys, join tables), indexes etc.



From Logical to Physical Design

- From the high-level design (concepts and ideas) to the low-level design (implementation)
- Logical modelling is for attributes used in the real world
 - You shouldn't be making attributes up in logical stages
 - Your natural keys uniquely identify entities in world
 - Your normalisation should be based on your logical modelling
- For people who have done some database design before, one of the first things people want to do in normalisation is start adding ID fields
 - But things like these these come later, on implementation, not at modelling
 - At the logical stage, you should only be thinking about the data that actually exists in the world

What can go wrong?

- Bunging an ID to uniquely identify a module
- But now you have to think about the data
 - We've got two different lecturers for the same module
- Breaks relationships in a database design
 - "Just link everything together by IDs"
- You're basically just linking spreadsheets together...

Module	Name	School	Module Leader	Lecturer
COMP1203	Computer Systems I	ECS	Kirk Martinez	Andrew Brown
COMP1203	Computer Systems I	ECS	Kirk Martinez	John Carter
COMP1204	Data Management	ECS	Oliver Bills	Danesh Tarapore
COMP1204	Data Management	ECS	Oliver Bills	George Konstantinidis
COMP3201	Cybersecurity	ECS	Ed Zaluska	Oliver Bills
COMP3210	Advanced Computer Networks	ECS	Kirk Martinez	Oliver Bills
ELEC1203	Mechanics	ECS	Igor Golosnov	Christopher Freeman
ELEC1207	Electronic Systems	ECS	Nick Harris	Paul Lewin
ELEC1207	Electronic Systems	ECS	Nick Harris	Rob Maunder
FEEG1003	ThermoFluids	FEE	John Shrimpton	
PHYS1011	Waves, Light and Quanta	Physics	Andrew Akeroyd	David Smith
PHYS1013	Energy and Matter	Physics	Pierre Thibault	Pasquale di Bari

The classic example

Primary key (Suit, Card)

Suit	Card
Hearts	Ace
Hearts	2
Hearts	3
Hearts	4
Hearts	5
Hearts	6
Hearts	7
Hearts	8
Hearts	9
Hearts	10
Hearts	Jack
Hearts	Queen
Hearts	King

Primary key (ID)

ID	Suit	Card
1	Hearts	Ace
2	Hearts	2
3	Hearts	3
4	Hearts	4
5	Hearts	5
6	Hearts	6
7	Hearts	7
8	Hearts	8
9	Hearts	9
10	Hearts	10
11	Hearts	Jack
12	Hearts	Queen
13	Hearts	King



Surrogate Keys: ID Fields

- Automatic ID fields are what we term surrogate keys: They have **no** business meaning
- Surrogate keys are attributes created and maintained **by the system** to aid in uniquely identifying an instance of an entity – they do not occur in the real world
- They are used because natural keys aren't generally stable, consistent or efficient to use in a database system
- They stand in place of the natural keys for technical purposes only, and are added as part of physical design – after the logical modelling and normalisation
- But you don't need to use them everywhere – some developers want an ID in every table...

ID	Module	Name
1	COMP1203	Computer Systems I
2	COMP1204	Data Management
3	COMP3201	Cybersecurity
4	COMP3210	Advanced Computer Networks
5	ELEC1203	Mechanics
6	ELEC1207	Electronic Systems
7	FEEG1003	ThermoFluids
8	PHYS1011	Waves, Light and Quanta
9	PHYS1013	Energy and Matter

ID	Lecturer
1	Andrew Brown
2	Christopher Freeman
3	Danesh Tarapore
4	David Smith
5	George Konstantinidis
6	John Carter
7	Oliver Bills
8	Pasquale di Bari
9	Paul Lewin
10	Rob Maunder

Module_ID	Lecturer_ID
1	1
1	6
2	3
2	5
3	7
4	7
5	2
6	9
6	10
8	4
9	8

More keys


- **Primary key:** a key with attributes that are not allowed to be **Null**
- **Foreign key:** an attribute of one relation that references an attribute (primary key) of another relation
 - Employee(ID,Name,**DeptNo**)
 - Department(ID,Name)

Conceptual Data Models


- Include important entities and the relationship between them.
- Do not specify attributes.
- Do not specify primary keys.
- Used as the foundation for logical data models.

Logical Data Model

- Include all entities and relationships between them.
- Specify attributes for each entity.
- Specify primary key for each entity.
- Specify foreign keys, which identify the relationship between different entities.
- Involve normalization




Logical data models



Directionality indicates a constraint

Physical Data Model

- Specify all tables and columns.
- Include foreign keys to identify relationships between tables.
- May include normalization, depending on user requirements.
- May be significantly different from the logical data model.
- Will differ depending on which DBMS (database management system) is used.



Entity


- An Object we identify in our system
- Entities have **attributes** (e.g., Employee has name/dept)
- Some of these attributes may *functionally determine others* (see previous lectures)
- So how do we visually express entities/relationships?

Conceptual Model of Entities



- **Strong entities** exist independently from other entity types. They always possess a key.
- **Weak entities** depend on some other entity type (e.g., Representing Employees in Projects using an EmployeeProject entity)
- **Associative entities** are entities that associate the instances of one or more entity types (e.g., Representing Matches played by Players from a certain Team)

Relationships



- **Relationships** are meaningful associations between or among entities. A relationship provides useful information that **could not be discerned** with just the entity types.
- **Weak relationships:** connections that exist between a weak entity type and its owner (e.g., EmployeeProject)

Attributes




- **Attributes** : characteristics of either an entity, a many-to-many relationship, or a one-to-one relationship.
- **Multivalued attributes**: take on more than one value.
- **Derived attributes**: value can be calculated from related attribute values.

Turning Conceptual into Logical Model

- Primary Keys: underlined
- Foreign Keys: underlined
- Add attributes

Example: Conceptual to Logical



Physical ERDs

- Table structures,
 - column name, column data type, column constraints,
 - primary key, foreign key
 - relationships between tables

Example: Portfolio of Research Projects

- **Entities:**

- Employees
- Research Projects

- **Relationships:**

- Employees are associated with one dept.
- Projects have multiple Employees associated, with different roles
- Team leads

Types of relationships

- Cardinality:
 - **one-to-one** (e.g., an employee has one address)
 - **one-to-many** (e.g., an employee may be involved in multiple projects)
 - **Many-to-many** (e.g., a number of employees may be involved on many projects)
- Identifying/non-Identifying: when foreign key is/not part of the primary key of a child table.
- Mandatory/Optional
 - Employee must have a department
 - Department may have one or more employees
 - An EmployeeProject entity must have 1 employee and 1 project

Relationships (Crow's Foot Notation)



One



Many



One (and only one)



Zero or one



One or many



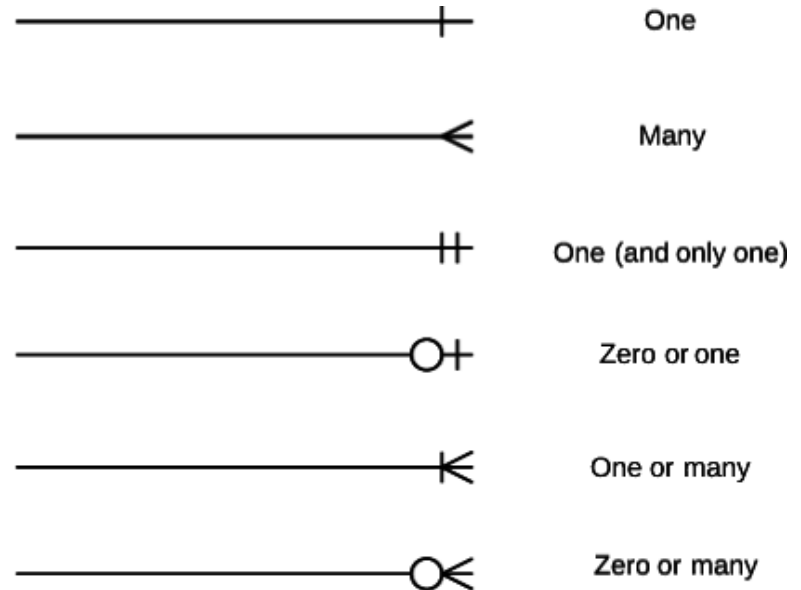
Zero or many

Cardinality and Modality

- **Cardinality** and **Modality** are the indicators of the business rules around a relationship.
- **Cardinality** refers to the **maximum number of times** an instance in one entity can be associated with instances in the related entity.
- **Modality** refers to the **minimum number of times** an instance in one entity can be associated with an instance in the related entity.

Modality and Cardinality

- Modality: min
- Cardinality: max



Representing Entities

Employee	
PK	<u>ID</u>
	Name
FK	DeptNo

Dept	
PK	<u>ID</u>
	DeptName

Project	
PK	<u>ID</u>
	Name
FK	DeptNo


EmpProject	
FK	EmpID
FK	ProjID
	Position

Indicate Key type
(underline PKs)

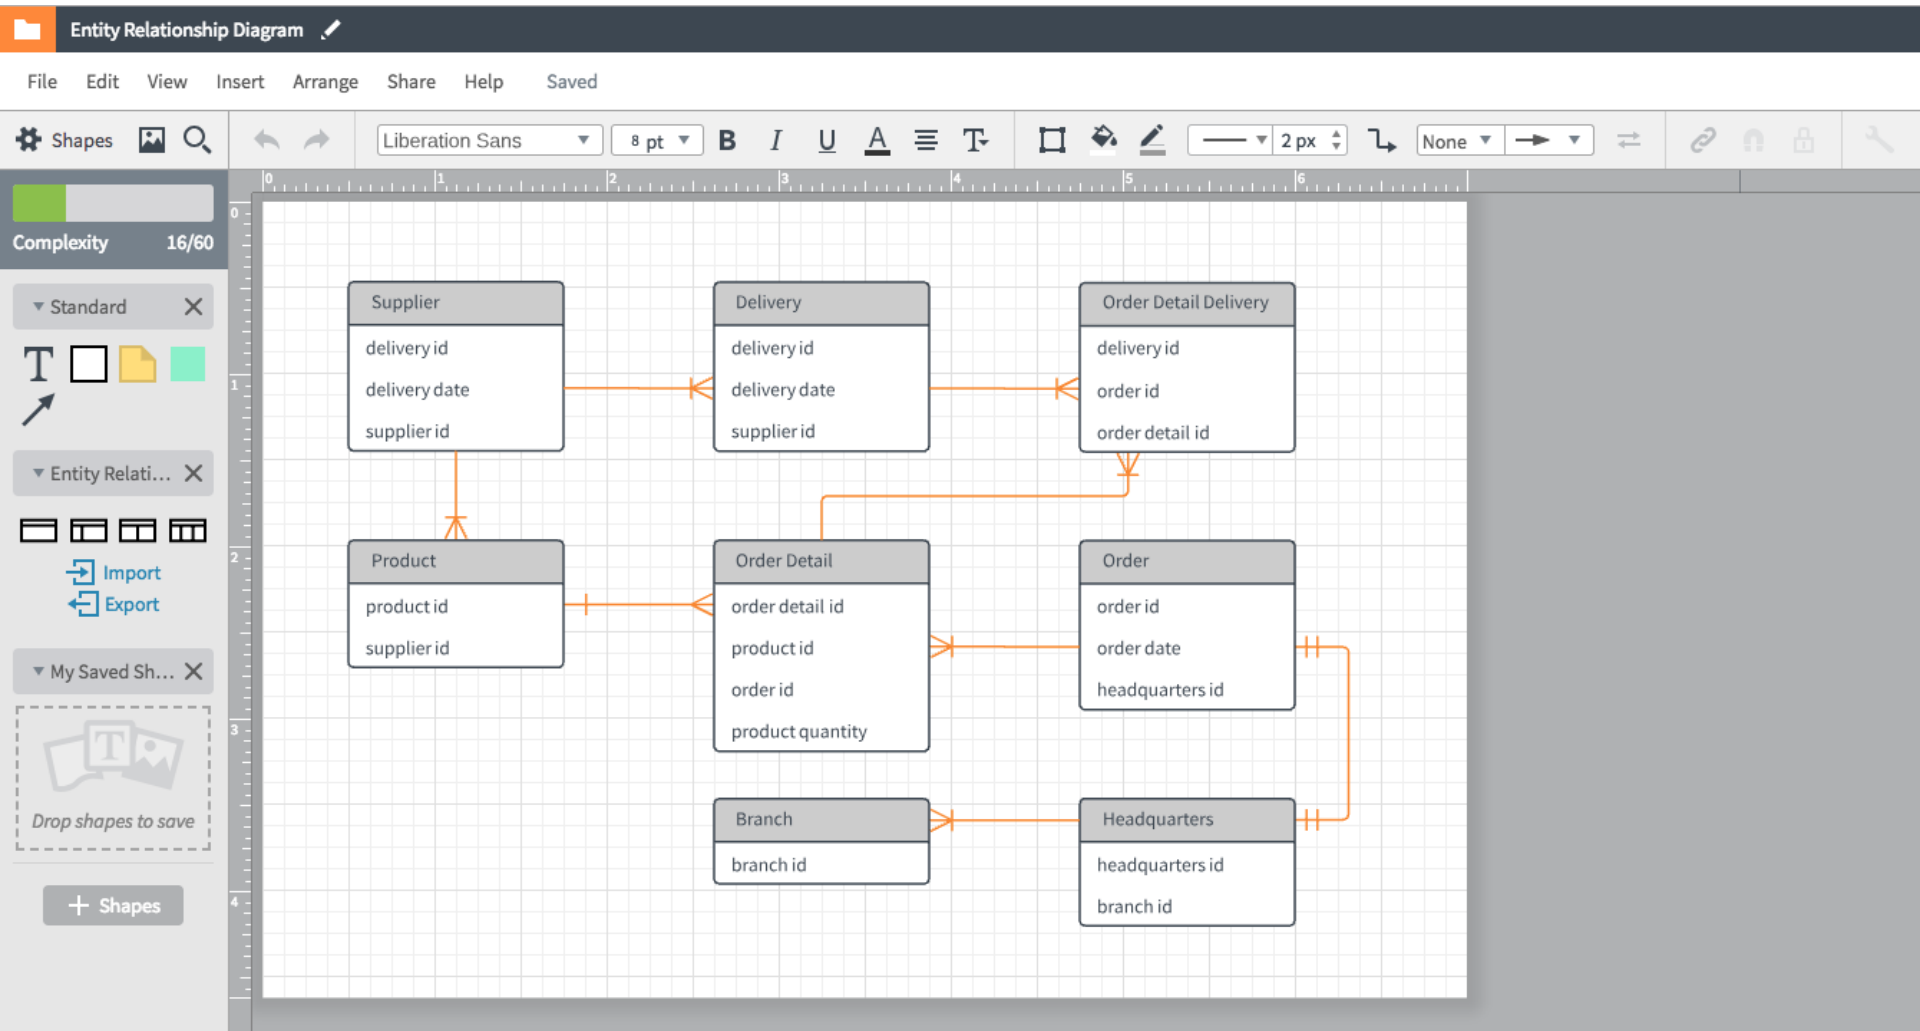
bold for non-null

Drawn using
LucidChart

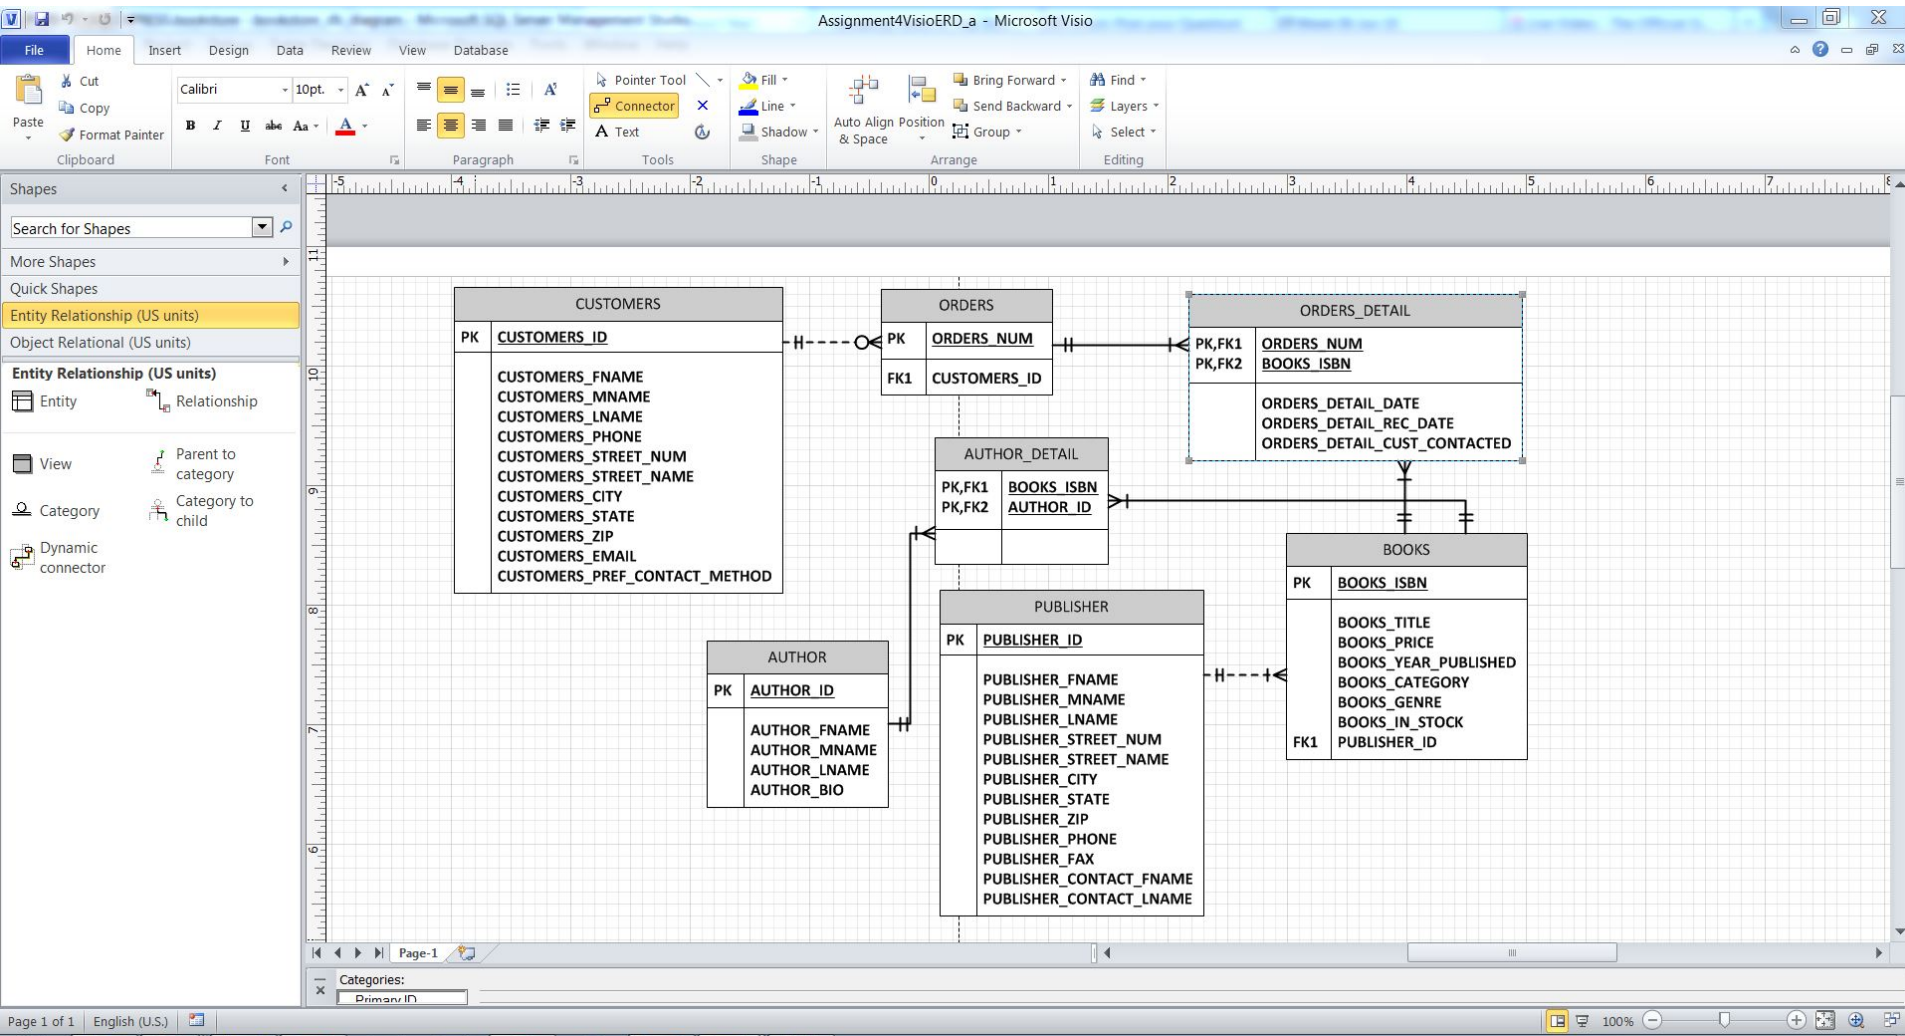
An ERD for Research Projects



Making ERDs



Making ERDs



Visio is available on vdi.soton.ac.uk, lab machines and from <http://ecs.gg/imagine>



Product Search



Announcements

Windows 10 — Important! 2017-08-16

Popular

Operating Systems

Developer Tools

Servers

Applications

Training

All



Microsoft Azure for Students
Starter



Windows 10



Visual Studio Enterprise
2017



Visual Studio Community
2017



Visual Studio for Mac



SQL Server 2017 Developer



SQL Server 2017 Enterprise



SQL Server 2017 Standard



SQL Server 2017 Web



Windows Server 2016



Access 2016



Project 2016



Visio 2016



Windows 8.1



Access 2013



Project 2013



Visio 2013

From Modelling to SQL

What is SQL



- **SQL: Structured query language**
- Specifies a **Data Definition Language (DDL)**
 - Tables and views (virtual tables).
 - Convert a data model to a (physical) database
- Specifies a **Data manipulation (DML)**
 - Programmatic data manipulation
 - Declarative (desired result)
 - INSERT, DELETE, UPDATE or retrieve (SELECT) data.
- Enforces Data integrity:
 - Referential integrity
 - Transactions
 - Checks keys for consistency
- Access control: security
- Data sharing: by concurrent users



SQL vs Programming Languages

- Restricted language : about 30 sub statements
- Restricted number of operations -> consistency across systems
- Hides low-level data operations
- Focus on programming data manipulation rather than control loops

SQL

- **Data definition** : define tables and views.
- **Data query** : extract data, add data and delete data.
- **Administration** : grant permissions to users to perform operations on our database.